

Exp C: does the model’s pre-decoding entropy track its own decoder β ?

May 13, 2026

1 Question

The project proposal hypothesizes (`decoding-modeling.tex`): *language models form a belief about their own decoding parameters by conditioning on the text they have generated*. This experiment tests that hypothesis directly: holding the prompt fixed, vary β_{dec} during self-generation and ask whether the model’s pre-decoding entropy $H(\text{softmax}(\ell_t))$ at late steps reflects the decoder β_{dec} in use.

2 Setup

Qwen-2.5-3B base, 32 prompts (`src/decoding_decoding/prompts.py`), each truncated to its first 10 tokens. Generate 150 tokens per prompt at three β_{dec} settings:

$$\beta_{\text{dec}} \in \{0.5 (T = 2.0), 1.0 (T = 1.0), 2.0 (T = 0.5)\}.$$

At each step, record $H(\text{softmax}(\ell_t))$ *before* the decoder applies β_{dec} . Late-window aggregation: per-prompt mean H over steps 80–149.

3 First-look result

Table 1: Late-window (steps 80–149) pre-decoding entropy per β_{dec} , with per-prompt paired difference.

condition	late-window H mean (nats) $\pm$ SE	paired diff vs $\beta = 2.0$
$\beta_{\text{dec}} = 0.5$	9.216 ± 0.052	$+8.045 \pm 0.115$ ($t = +69.7$, $n_+ = 32/32$)
$\beta_{\text{dec}} = 1.0$	2.168 ± 0.186	$+1.00$
$\beta_{\text{dec}} = 2.0$	1.172 ± 0.109	—

The late-window entropy moves ≈ 8 nats (factor ≈ 3000 in effective vocabulary) between $\beta_{\text{dec}} = 0.5$ and $\beta_{\text{dec}} = 2.0$. Every one of 32 prompts orders the same way.

Three observations before we believe the headline.

1. Natural high-entropy human writing (WritingPrompts/natural) sits at $H \approx 2.85$ nats — nowhere near the 9.2 nats seen here. So this is not the model ”reading off high-T-like content”; this is a regime that only off-distribution self-sampling reaches.




Figure 1: Median $H(t)$ per condition, IQR shaded. Dashed grey lines mark mean entropies of real natural text from the earlier natural_text experiment (Wikipedia ≈ 1.98 , WritingPrompts ≈ 2.85 nats).

2. Inspection of $\beta = 0.5$ generations (Figure 2) shows pure token salad: code keywords, foreign-language fragments, URL slugs, random named entities. $\beta = 1.0$ generations are coherent English. $\beta = 2.0$ generations are also coherent English, more compressed toward common factual phrasing.
3. The entropy trajectory at $\beta = 0.5$ climbs monotonically from ≈ 2 at step 0 to ≈ 9 at step 80: classic compounding off-distribution drift, where each random token makes the next context less parseable and the next sample yet more random.

These three points raise an obvious alternative: the 8-nat effect is *garbage-in, garbage-out* — the model is uncertain because its context is incoherent, not because it has formed a “belief” that the decoder is set to high T .

4 Swap-prefix control

To disambiguate, we hold the early prefix’s generation regime fixed and *switch* the sampler β at step 80. If the model truly has a belief about β_{dec} that updates on recent tokens, then switching to $\beta = 1.0$ at step 80 should pull late-window H down to the $\beta_{\text{dec}} = 1.0$ baseline (≈ 2 nats) within ~ 20 tokens. Five conditions, same 32 prompts:

- pre05_post05 — $\beta = 0.5$ throughout (no switch, control).
- pre05_post10 — $\beta = 0.5 \rightarrow 1.0$ at step 80. **This is the decisive comparison.**
- pre05_post20 — $\beta = 0.5 \rightarrow 2.0$ at step 80.
- pre20_post05 — $\beta = 2.0 \rightarrow 0.5$ at step 80 (mirror).

```

 $\beta_{\text{dec}} = 0.5$  (T = 2.00):
«Marie Curie was a Polish-born physicist and chemist whose pioneering research on» a very اءءءء spent $ ;\n Crefit0reement.ToInt_Begin(' About website flood
entirety CommonTouchGMT not[] institution going Bing0000 [][] []stral_debugignon>;\n \n inundlletcherYZ merchant directory genis_Bool Jasmine ast_operations BSP
PV0ID[] []As.HasPrefix_media BUG),\n 0 extr[] []Fal\ent[] outward

 $\beta_{\text{dec}} = 1.0$  (T = 1.00):
«Marie Curie was a Polish-born physicist and chemist whose pioneering research on» chemist who made significant contributions to the study of radioactivity.
Although research is ongoing, it is said that at least 60 percent of the total alpha decay energy of all the potassium 40 in the human body (there are 15 Bq/Kg
potassium 40),

 $\beta_{\text{dec}} = 2.0$  (T = 0.50):
«Marie Curie was a Polish-born physicist and chemist whose pioneering research on» chemist who was the first woman to win a Nobel Prize, and the first person
to win two Nobel Prizes, one in physics and the other in chemistry. She was born on November 7, 1867 in Warsaw, Poland. Her father, Wladyslaw 0ls

```

Figure 2: First 60 sampled tokens for prompt 0 (Marie Curie) at each β_{dec} . $\beta = 0.5$ generation degrades to token salad within ≈ 20 steps.

- **pre20_post20** — $\beta = 2.0$ throughout (control).

Table 2: Pre-switch median H and post-switch median H at step offsets from the switch.

condition	$\beta_{\text{pre}} \rightarrow \beta_{\text{post}}$	pre-med	+0	+5	+10	+20	+50
pre05_post05	0.5 $\rightarrow$ 0.5 (no switch)	8.42	9.26	9.21	9.49	9.66	9.47
pre05_post10	0.5 $\rightarrow$ 1.0	8.56	9.44	9.10	9.10	8.90	8.92
pre05_post20	0.5 $\rightarrow$ 2.0	8.57	9.53	6.83	4.33	1.68	0.33
pre20_post05	2.0 $\rightarrow$ 0.5	1.21	0.69	5.20	6.46	8.01	9.06
pre20_post20	2.0 $\rightarrow$ 2.0 (no switch)	1.29	0.64	0.58	1.05	0.30	0.39

What this shows.

- **pre05_post10**: switching the sampler from $\beta = 0.5$ to $\beta = 1.0$ leaves median H essentially unchanged (8.56 $\rightarrow$ 8.92 at step +50). The model does *not* update its predictions to match the new sampling regime over even 50 tokens of new evidence.
- **pre05_post20**: switching to $\beta = 2.0$ *does* recover coherence over ~ 20 tokens. But this recovery is explained by the mechanism of $\beta = 2$ sampling itself — $\beta = 2$ forces a top-1 selection from $\text{softmax}(\beta \cdot \ell_t)$, which under a flat ℓ_t still picks the highest-probability token (typically a common safe continuation like punctuation, common articles, frequent function words); those tokens then *re-anchor* the context to in-distribution English. This is not the model "updating its belief"; it is the sampler injecting in-distribution tokens that gradually overwrite the bad prefix.

The decisive comparison **pre05_post10** falsifies a strong form of the "decoder-belief" framing. If model-belief-about- β were a fast posterior update on recent sampled tokens, switching the sampler should be visible within a handful of tokens. It is not.




Figure 3: Median pre-decoding H at offsets $+0, +5, +10, +20, +50$ after the sampler switch. Critical observation: *the orange line (pre05_post10) does not fall* — the model does not recover coherence when the sampler switches to $\beta = 1.0$.

5 What this does *not* say

The project’s broader claim — that the model forms beliefs about its own decoder — is *not* falsified by this experiment. Prior work in this repo gives direct positive evidence:

- **Blacklist self-prediction** (results/blacklist_summary.md, run 2026-05-10): when token ranks 1–5 are banned at sampling, the model’s softmax probability on the banned token drops between early and late positions, with cluster-bootstrap CIs that do not cross zero. The rank-1 ban (‘ the’) shows $\Delta p_{\text{late-early}} = -0.018$ and a paired late treatment-vs-control difference of -0.016 (CI $[-0.019, -0.013]$). The model learns its own blacklist from the absent-token signature.
- **Top- K self-prediction** (results/summary_L2048.md, run 2026-05-09): under top- k sampling, the cliff at rank K in the model’s predicted-rank distribution exceeds the corresponding cliff under unrestricted sampling. Treatment minus control at rank K is $+0.475$ for $k = 3$ and $+0.170$ for $k = 5$ (CIs tight, all clear of zero). Within-treatment concentration grows between early and late positions for $k = 1, 3, 5$ and *shrinks* for $k = \infty$: this is not generic drift.

Together with the earlier natural-text WP offset ($\Delta \log \hat{\beta} \approx -0.24$, $t = -3.8$ on real creative writing, natural_text/findings.pdf), these say the model does adjust its predictions in response to which tokens were sampled in its context.

6 What this experiment *does* say

Exp C tested a specific strong instantiation of decoder-belief: the model’s *pre-decoding entropy* as a fast posterior readout that snaps to the new sampler β within a few dozen tokens of mid-generation sampler change. That specific form fails. The swap-prefix control shows the model does not pull its entropy back to the $\beta = 1$ baseline over 50 tokens of new evidence under a different sampler.

So the belief-update mechanism is not a fast posterior read off per-token entropy. It plausibly operates on slower, more aggregate features — rank statistics, support shape, blacklist signatures — exactly the signals that the blacklist and top- k experiments isolate.

The remaining ambiguity in Exp C: the 8-nat raw effect is a mix of genuine decoder-belief signal (consistent with blacklist/top- k findings) and garbage-in/garbage-out off-distribution drift from $\beta = 0.5$ sampling. The two cannot be separated cleanly in this experimental design because $\beta = 0.5$ decoheres the prefix faster than any belief update could compensate.

7 Bottom line

- The 8-nat entropy spread between $\beta_{\text{dec}} = 0.5$ and $\beta_{\text{dec}} = 2.0$ self-generation is a mix of (i) real decoder-belief signal already established by the blacklist and top- k experiments and (ii) garbage-in/garbage-out off-distribution drift. This experiment cannot separate them.
- The swap-prefix control rules out the specific strong instantiation in which pre-decoding entropy is a fast posterior readout that snaps to the new sampler β within ~ 50 tokens.
- The broader framing — that the model forms beliefs about its decoder by conditioning on its context — remains supported by the blacklist and top- k experiments. Exp C constrains the form of that belief update rather than refuting its existence.
- A clean follow-up isolating decoder-belief from content quality would condition on *natural* text of varying entropy (rather than self-generated extreme- β samples) and ask the same questions about $H(t)$ and rank statistics.